

Installing k8s on Ubuntu 20.04 or ubuntu 22.04

1) Update your system.

sudo apt update

2) Add repository for docker.

sudo mkdir -p /etc/apt/keyrings

curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o /etc/apt/keyrings/docker.gpg

echo "deb [arch=\$(dpkg --print-architecture)
signed-by=/etc/apt/keyrings/docker.gpg]
https://download.docker.com/linux/ubuntu \
\$(lsb_release -cs) stable" | sudo tee /etc/apt/sources.list.d/docker.list >
/dev/null

3) Configure IP Address mapping in “/etc/hosts” file

Find IP's of your vm/systems by ‘ip a” command on enxxx interface

E.g

#####

#BELOW LINES NEED CUSTOMIZATION

#echo " " >> /etc/hosts
#echo "192.168.1.11 w1" >> /etc/hosts
#echo "192.168.1.10 cp" >> /etc/hosts

4) Set hostname of the systems

For control plane, on control plane node, execute following command:-

#hostnamectl set-hostname cp

5) On Worker node you will use following command.

#hostnamectl set-hostname w1

6) Load the Modules needed, create respective files as given below

modprobe br_netfilter
modprobe overlay

cat << EOF | tee /etc/modules-load.d/k8s-modules.conf

br_netfilter
overlay
EOF

```
cat << EOF | tee /etc/sysctl.d/k8s.conf
net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1
net.ipv4.ip_forward = 1
EOF
```

7) Update system control.

```
sysctl --system
```

8) Install containerd runtime

```
apt-get update ; apt-get install -y containerd.io
```

```
mkdir -p /etc/containerd
```

```
containerd config default | tee /etc/containerd/config.toml
```

```
sed -i "s/SystemdCgroup = false/SystemdCgroup = true/g"
/etc/containerd/config.toml
```

```
systemctl restart containerd
```

10) Turn of swap, remove or comment (hash) the entries for swap in “/etc/fstab” file.

```
swapoff -a
```

11) Install prerequisites required, by adding the repositories for k8s

```
apt-get install -y apt-transport-https curl
```

```
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key add
```

```
apt-add-repository "deb http://apt.kubernetes.io/ kubernetes-xenial main"
```

12) Install k8s binaries/ executables

```
apt install -y kubeadm kubelet kubectl
```

13) Install bash completion for command auto completing commands via tab press.

```
# install bash-completion
sudo apt-get install bash-completion
```

```
# Add the completion script to your .bashrc file
echo 'source <(kubectl completion bash)' >> ~/.bashrc
```

```
# Apply changes
source ~/.bashrc
```

14) After installation, execute kubeadm init on control plane.

```
sudo kubeadm init --pod-network-cidr=10.244.0.0/16 --apiserver-advertise-address=< IP address of control-plane>
```

15) Step 14, command returns with success - then use following commands on worker node, which is given by kubeadm init command, below is an example..

E.g

```
kubeadm join 192.168.1.251:6443 --token awilnt.qskrvqbkbng1tzea --
discovery-token-ca-cert-hash
sha256:869387b8708df3a0d587c9670f568292bb1aafd4776ee8f1f4fa5096208
e5dc0
```

16) There will be following command also needed for using the cluster as regular user

```
mkdir -p $HOME/.kube
```

```
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

16) Confirm with command for cluster nodes joining

```
kubectl get nodes
```

17) Configure CNI - flannel, for example for pod networking.

```
kubectl apply -f
```

<https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml>

18) Confirm again by running 'kubectl get nodes' for "ready status"

E.g

```
kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
k8ctrl	Ready	control-plane	5d3h	v1.28.2
u200w1	Ready	<none>	5d3h	v1.28.2
u200w2	Ready	<none>	5d3h	v1.28.2

19) Create deployment and check, with following command

e.g

```
kubectl create deployment nginx --image=nginx
```

```
kubectl get deployments nginx
```

```
kubectl create service nodeport nginx --tcp=80:80
```

```
kubectl get svc
```

```
curl localhost:30658
```

Note: 30658 port will vary from system to system. Check and use what's configured for your setup.

If all above, your cluster is successfully configured.